

IS5113 – AI Ethics and Regulation

Course Outline

Instructor: Prof Matthew LEE

Email: matthew.lee@cityu.edu.hk

“Ethics is knowing the difference between what you have a right to do and what is right to do.”

— Potter Stewart

“There is a longer-term existential threat that will arise when we create digital beings that are more intelligent than ourselves... if they are created by companies motivated by short-term profits, our safety will not be the top priority.”

— Geoffrey Hinton

“AI is a fundamental risk to the existence of human civilization... If AI is a potential risk to the public, then there should be some regulatory body that oversees what companies are doing”

— Elon Musk

Course Description:

This course reviews general principles of ethical theories and models and explores their application to AI in business and societal contexts. The course covers various AI ethics theories and governance frameworks, examining the relevant ethical, legal, and societal challenges involved in AI development and deployment. Students will critically analyze issues such as bias, transparency, privacy, accountability, security, intellectual property, and global policy, while building competence to design compliant corporate AI ethical guidelines and regulatory frameworks aligned with business strategy. The course blends theory, case studies, and practical projects to prepare future business leaders with AI ethics skills and knowledge necessary for success in a business world increasingly driven by AI and related innovations.

By the end of the course, students will be able to:

1. Understand core ethical theories and frameworks (e.g., utilitarianism, deontology, rights-based ethics) applicable to AI development and deployment.
2. Understand key contemporary AI ethical issues (e.g. fairness, accountability, transparency, privacy) in the context of business AI
3. Understand the business value of ethical AI.
4. Evaluate AI's business and societal impacts and examine how AI ethics and governance frameworks may address the challenges involved.
5. Apply relevant theories and frameworks to design corporate AI governance mechanisms and policies aligned with business strategy and compliant with global regulations (e.g. GDPR, EU AI regulations, China's AI laws).
6. Engage in critical thinking and informed debates on the ethical challenges brought about by global business AI.

Assessment: (To pass the course, you need to pass BOTH the written examination and the continuous assessment)

Final Written Examination (50%)

Continuous assessment (50%)

- Participation (10%): Actively engage with class discussions, quizzes, assignments, and case studies


- Group Project (40%): (a) A research investigation on a current AI ethics topic; (b) Design an AI governance strategy (or policy) for a specific industry (e.g., retail, fintech), company (real or fictitious), or business function (e.g., marketing)

Weekly Seminar Outline (Tentative):


Week	Topic
1.	Seminar: Introduction to AI and Ethics
2.	Seminar: Ethical Theories, Practical Applications and Limitations
3.	Seminar: Bias and Fairness
4.	Seminar: Transparency and Explainability
5.	Seminar: Accountability and Liability
6.	Seminar: Privacy and Surveillance
7.	Seminar: Ethical Issues - AI in Marketing, Consumer Behavior, HR & Hiring
8.	Seminar: Ethical Issues - AI in Finance & Risk Management, Generative AI & Misinformation
9.	Seminar: Global AI Governance & Regulation
10.	Seminar: Environmental & Social Impacts
11.	Seminar: Emerging Risks and Case Studies
12.	Project Presentation
13.	Project Presentation

Reading Materials

AI Ethics by Mark Coeckelbergh. The MIT Press, 2020.



AI Ethics: A Textbook by Pauline Boddington. Springer Nature Singapore, 2024



What is AI (Artificial Intelligence), McKinsey & Company, April 3, 2024
<https://www.mckinsey.com/featured-insights/mckinsey-explainers/what-is-ai>

Note: Reading material other than the course textbooks will be required as the course progresses