

BIAS, FAIRNESS & DISCRIMINATION IN AI

INSTRUCTOR:

PROF. MATTHEW LEE

MATTHEW.LEE@CITYU.EDU.HK





AGENDA

- Justice, Fairness, Bias: an introduction
- AI Bias
- Mitigation Strategies

INTRODUCTION

- Justice, Fairness, and Bias:
- Philosophical, Psychological, and Business Perspectives



UNDERLYING IDEAS

- Defining justice, fairness, bias
- Philosophical foundations
- Psychological mechanisms
- Business perspectives

WHAT IS JUSTICE?

- Concerned with rightful allocation of benefits and burdens
- Rooted in moral philosophy and social order
- Varies across cultures and historical periods
- Questions of distribution, recognition, and participation



HISTORICAL CORNERSTONES

- Aristotle's distributive and corrective justice
- Plato's vision of harmony in the Republic
- Kant's categorical imperative and respect for persons
- Hobbes, Locke, Rawls on social contract foundations

*JUSTICE AS
FAIRNESS
(RAWLS)*

"Justice consists...in treating equals equally and unequal unequally and in giving each person his due."

Two basic principles: "The distribution of benefits and burden in a society is just if and only if:

- 1. Each person has an equal right to the most extensive basic liberties compatible with similar liberties for all
- 2. social and economic inequalities are arranged so that they are both
 - a. to the greatest benefits of the least advantaged person
 - b. attached to offices and position open to all under conditions of fair equality of opportunity.
- Principle 1 takes priority over principle 2 if they come into conflict. Within principle 2, 2b takes priority over 2a

1. PRINCIPLE OF EQUAL LIBERTY

- Each citizen basic liberties must be protected from invasion by others and must be equal to those of others
 - Basic liberties: right to vote freedom of speech and conscience, civil liberties, freedom from arbitrary arrest.
 - E.g. contract must be free of fraud and must be honored. Employees must render services justly contracted with employer

2A. DIFFERENCE PRINCIPLE

- A productive society will incorporate inequalities but takes steps to improve the position of the neediest member of society

2B. PRINCIPLE OF FAIR EQUALITY OF OPPORTUNITY

- Everyone should be given an equal opportunity to quality for the more privileged position in society's institutions.
- E.g. Job qualifications should not only be related to requirements of the job but also each person must have access to training and education needed to quality for the job remuneration would depend on person efforts abilities and contribution

*RAWLS'S
ORIGINAL
POSITION
AND VEIL OF
IGNORANCE*

- Imagining decisions behind a veil of ignorance
- Prioritizing the least advantaged (maximin rule)
- Balancing liberty with fair equality of opportunity

RAWLS METHOD

- Rawls theory provides us not only with a set of principles of justice but also a general method for an evaluating the adequacy of any moral principle in a fair way
- The method consists of determining what principles a group of rational self-interested persons would choose to live by if they knew they would live in a society governed by those principles, but they did not yet know what each of them would turn out to be like in that society."

The image shows a chalkboard with handwritten mathematical derivations. On the left, a graph of a function $y = g(x)$ is shown with a point $x+h$ on the x-axis. A secant line connects the point $(x, g(x))$ to the point $(x+h, g(x+h))$. A tangent line is drawn at the point $(x, g(x))$. The labels "Secant Lines" and "Tangent Line" are written next to their respective lines. On the right, the derivative $f'(x)$ is defined as the limit of the difference quotient:

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Below this, the derivative of $f(x) = x^2$ is calculated:

$$\begin{aligned} f'(x) &= \lim_{h \rightarrow 0} \frac{(x+h)^2 - x^2}{h} \\ &= \lim_{h \rightarrow 0} \frac{x^2 + 2xh + h^2 - x^2}{h} \\ &= \lim_{h \rightarrow 0} \frac{2xh + h^2}{h} \\ &= \lim_{h \rightarrow 0} h(2x + h) \end{aligned}$$

FAIR PRINCIPLE

- 'A principle is a morally justified principle of justice if and only if the principle would be acceptable to a group of rational self-interested persons who know they will live in a society governed by the principles they accept but who do not know what sex, race, abilities, religion, interests, social position, income, or other particular characteristics each of them will possess in that future society

VEIL OF IGNORANCE

Original position "An imaginary meeting of rational self-interested persons who must choose the principles of justice by which their society will be governed."

Veil of ignorance "The requirement that persons in the original position must not know particulars about themselves which might bias their choices such as their sex, religion, race, income, etc."

INLINE WITH KANTIAN REQUIREMENTS

Reversibility (parties choose principles they apply to themselves)

Reversalizable (principles must apply equally to everyone)

Treating people as ends (each party has equal say in choice of principles)

Rawls claim that parties to the original position would in fact choose these principles of justice: principle of equal liberty, difference principle and principle of fair equality and opportunity.

EQUALITY

VS.

EQUITY



EQUALITY = SAMENESS
GIVING EVERYONE THE SAME THING



EQUITY = FAIRNESS
ACCESS TO SAME OPPORTUNITIES

MAJOR THEORIES OF DISTRIBUTIVE JUSTICE

- Utilitarianism: greatest good for the greatest number
- Egalitarianism: equal access to resources
- Libertarianism: entitlement and property rights
- Rawls's theory of justice as fairness and the difference principle

PROCEDURAL VS. DISTRIBUTIVE VS. INTERACTIONAL JUSTICE

- Distributive: fairness of outcomes
 - **Distributive Justice:** involves the fair distribution of benefits and burdens.
 - **Retributive Justice:** concerns the "Fairness when blaming or punishing person for doing wrong."
 - **Compensatory justice:** concerns "the fairness when restoring to a person what the person lost when he or she was wronged by someone else."
- Procedural: fairness of decision-making processes
- Interactional: quality of interpersonal treatment
- Philosophical overlaps and distinctions

THEORIES OF DISTRIBUTIVE JUSTICE

1. Justice as Equality: Egalitarianism

2. Justice Based on Contribution: Capitalist Justice

3. Justice Based on Needs and Abilities: Socialism

4. Justice as Freedom: Libertarianism

5. Justice as Fairness: John Rawls

THE DOMINANT VIEW OF FAIRNESS (RAWLS)

- The principle of distributive justice simply states:
 - Individuals who are similar in all respects relevant to the kind of treatment in question should be given similar benefits and burdens even if they are dissimilar in other irrelevant respects and individuals who are dissimilar in a relevant respect ought to be treated dissimilarly in proportion to their dissimilarity.
 - *E.g. Sales persons A and B have the same job positions but at the end the year. A sold more than B therefore A gets the merit increase or promotion. Distributive justice principle purely formal. Principle does not indicate "relevant respects" for similar or dissimilar treatment. Material principle of justice :different theories of distributive justice more fundamental "relevant respects" than just "first -come-first-serve principle" or "priority for senior citizens."

PHILOSOPHICAL BIAS

- Are moral principles universal or culturally bound?
- Self-interest vs. altruism in Kantian and utilitarian thought
- Overly simplistic premises when applied to real human behavior
- Philosophical caution on prescriptive vs. descriptive claims

PSYCHOLOGICAL PERSPECTIVE

- Philosophical ideals meet real-world minds
- How do individuals perceive and enact fairness?
- Introduction to cognitive and motivational influence

DEFINING FAIRNESS IN PSYCHOLOGY

- Subjective experience of equity, equality, and reciprocity
- Fairness as a basic social heuristic
- Cultural norms and individual differences
- Link between perceived fairness and well-being

HEURISTICS IN FAIRNESS JUDGMENTS



- Mental shortcuts for quick evaluations



- Equality heuristic: split resources evenly



- Equity heuristic: proportional to contribution



- Procedural heuristic: process matters more than outcome

*COGNITIVE
BIASES
AFFECTING
FAIRNESS*

Biases are unconscious and automatic processes designed to make decision-making quicker and more efficient. Cognitive biases can be caused by many things, such as heuristics (mental shortcuts), social pressures, and emotions. Broadly speaking, bias is a tendency to lean in favor of or against a person, group, idea, or thing, usually in an unfair way

"people use mental approximations to understand an uncertain world. As a result, we make certain types of errors in judgment." (Tersky)

- Self-serving bias: interpreting rules to one's advantage
- Confirmation bias: seeking evidence that supports one's view
- Loss aversion bias: negative outcomes weigh more than positive outcome
- Anchoring bias: first impression rules!
- Framing bias: it doesn't matter what you say, it is about how you say it

The list goes on




Self serving bias

Favouring decisions that enhance self-esteem. This results in attributing positive events to oneself and conversely negative events as blame on oneself. Within innovation workshops this can mean that decisions made can be loaded with personal agendas rather than customer and business logic for the company.




Confirmation bias

We believe what we want to believe by favouring information that confirms pre-existing beliefs or preconceptions. This results in looking for creative solutions that confirm our beliefs rather than challenge them.



Loss aversion bias

Once a decision has been made, sticking to it rather than taking risks due to the fear of losing what you gained in starting something and wishing to see it finished. We also attach more value to something once we have made an emotion investment in it. As a consequence of effort, time and energy put into creative thinking, team members can become biased and become emotionally attached to their outcomes. To remedy this, the 11th commandment: “thou shalt not fall in love with thy solutions”.



The value function that passes through the reference point is s-shaped and **asymmetrical**. The value function is steeper for losses than gains indicating that losses outweigh gains.

Prospect theory



Anchoring bias

Being influenced by information that is already known or that is first shown. This causes pre-loaded and determined tunnel vision and influences final decision making. We deliberately manipulate team members' minds by 'pre-loading' them one of our warm-up exercises to demonstrate this bias at play. The impact is highly-significant on creative thinking and outcomes.



Framing bias

Being influenced by the way in which information is presented rather than the information itself. We see this one all the time particularly when developing prototypes for pitching as well as in presenting polished slides. People will avoid risk if presented well and seek risk if presented poorly meaning that decision making logic can easily be skewed.

EMOTION AND FAIRNESS

- Moral outrage and retributive impulses
- Empathy and the drive to help the disadvantaged
- Anger at unfair treatment fuels social justice movements
- Regulatory focus: promotion vs. prevention orientation

DISTRIBUTIONAL FAIRNESS EXPERIMENTS

- Ultimatum game: responders reject low offers

- [The Ultimatum Game Psychology](#)

- Dictator game: altruism vs. selfishness in giving

- Allais and reflection effects in resource allocation

- Implications for policymakers

PROCEDURAL FAIRNESS EFFECT



- Voice: opportunity to present one's case



- Neutrality: unbiased decision-makers



- Trust: transparency and integrity of process



- Respect: dignified interpersonal treatment

INTERACTIONAL FAIRNESS

- Informational justice: explanations and communication
- Interpersonal justice: courtesy, respect, and empathy
- Role of leader behavior in shaping perceptions
- Impact on team cohesion and trust

MEASUREMENT OF FAIRNESS PERCEPTIONS



Surveys: organizational justice scales



tasks

Behavioral indices: resource allocation




Physiological markers: stress responses to unfairness



Big data approaches to tracking equity

BUSINESS PERSPECTIVE



- From individual minds to organizational structures



- Why fairness and justice matter in business settings



- Link between justice perceptions and performance

ORGANIZATIONAL JUSTICE FRAMEWORK

- Distributive justice: fair pay, benefits, and promotions
- Procedural justice: transparent policies and procedures
- Interactional justice: respectful HR practices
- Informational justice: clear, honest communication

IMPACTS ON EMPLOYEE BEHAVIOR



- Job satisfaction and organizational commitment



- Trust in leadership and willingness to cooperate



- Turnover intentions and absenteeism



- Creativity and willingness to go the extra mile

BIAS IN HIRING AND PROMOTION



- Unconscious stereotypes in talent identification



- Structured interviews and blind resume screening



- Performance appraisals and calibration meetings



- Inclusive leadership practices

FAIR COMPENSATION SYSTEMS



- Pay equity analyses across demographics



- Merit-based vs. seniority-based raises



- Transparent bonus criteria and variable pay



- Role of external benchmarks

DESIGNING FAIR PROCESSES



- Cross-functional decision panels



- Employee voice mechanisms and appeals processes



- Rotating chair for grievance committees



- Technology to ensure consistency and audit trails

CASE STUDY: IMPLEMENTING PROCEDURAL JUSTICE



- Scenario: roll-out of a new performance rating system



- Steps taken for transparency, employee input, and training



- Results: increased perceived fairness and productivity



- Lessons learned and best practices

INTEGRATING PERSPECTIVES




- Philosophical ideals guide organizational values



- Psychological insights shape policy design



- Business applications translate theory into practice



- Continuous feedback loops and evidence-based refinement

SUMMARY

- Justice spans moral theory, human cognition, and organizational design
- Fairness is multifaceted: outcomes, processes, and interactions
- Bias emerges at every level but can be mitigated by design
- Integrative approach yields sustainable, ethical, and effective system

The background features several thin, light-colored lines intersecting across the frame. A large, light-colored triangle is positioned on the right side, pointing towards the top right corner.

AI BIASES

AGENDA

- Define AI bias
- Identify sources of bias across the AI lifecycle
- Understand why fairness matters in business strategy
- Explore real-world examples and ethical dilemmas

WHY THIS MATTERS FOR BUSINESS LEADERS

- AI decisions affect hiring, lending, pricing, and more
- Biased systems can lead to:
- Legal action (e.g., discrimination lawsuits)
- Reputational damage (e.g., viral backlash)
- Financial loss (e.g., customer churn)
- Fairness builds trust, loyalty, and long-term value

WHAT IS AI BIAS

- AI bias, also called [machine learning](#) bias or algorithm bias, refers to the occurrence of [biased results](#) due to human biases that skew the original training data or AI algorithm—leading to distorted outputs and potentially harmful outcomes. (source: IBM).
- These biases can lead to distorted outputs that disproportionately affect certain groups, perpetuating discrimination, stereotypes, or systemic inequalities.

AI IN PRACTICE (RECAP)

- AI = systems that mimic human cognition
- Includes:
- Machine Learning (ML)
- Natural Language Processing (NLP)
- Computer Vision
- Learns patterns from data to make predictions or decisions

Example:

Netflix recommending shows, banks approving loans, HR filtering resume

Core Concepts of AI: Machine Learning and Beyond

Data

Essential for training algorithms and improving performance.

Training

Involves teaching models through labeled datasets to improve accuracy.

Automation

Streamlines repetitive tasks to increase operational efficiency.

Scaling

Allows deployment of AI solutions to handle increasing data volume.

Algorithms

Provide the framework for processing and analyzing data efficiently.

Evaluation

Uses metrics to assess the effectiveness of AI models.

Integration

Combines AI solutions with existing systems for enhanced capabilities.

Ethics

Ensures responsible use of AI technology in real-world applications.



TYPICAL AI SYSTEMS IN BUSINESS

- Predictive Analytics: Credit scoring, fraud detection
- Recommendation Engines: E-commerce, media
- Chatbots & Virtual Assistants: Customer service
- Automated Decision Systems: Hiring, insurance approvals

BIAS VS. FAIRNESS

- Bias: Systematic deviation from expected outcomes, often disadvantaging certain groups
- Fairness: Equitable treatment across individuals or groups
- Tension: Improving fairness may reduce accuracy or efficiency



DEFINITION AND IMPORTANCE OF BIAS

Nature of AI Bias

Bias in AI arises from systematic errors in data collection, algorithm design, or decision-making processes.

Impact on Fairness

Bias leads to unfair outcomes in sensitive areas like hiring, lending, and healthcare decisions.

Sources of Bias

Bias stems from historical inequalities in training data and underrepresentation of certain groups.

Need for Mitigation

Addressing bias is essential to build trustworthy AI systems serving all populations equitably.

TYPES OF BIAS IN AI SYSTEMS

Historical Bias

Historical bias results from training data reflecting past inequalities, perpetuating discrimination in AI systems.

Representation Bias


Representation bias occurs when certain groups are underrepresented in datasets, causing skewed AI outcomes.

Measurement Bias

Measurement bias arises from inaccurate quantification of features or outcomes, misrepresenting reality in AI models.

Aggregation and Evaluation Bias

Aggregation bias ignores group differences applying models uniformly; evaluation bias uses unsuitable benchmarks for testing.



HISTORICAL CONTEXT OF BIAS IN TECH

- Redlining Algorithms: Mortgage denials based on zip codes
- Facial Recognition Failures: Misidentification of Black and Asian faces
- Resume Screening Bias: Penalizing “female-coded” terms

Bias in tech isn't new—it's just more powerful now. AI can scale discrimination faster than humans ever could.

CASE STUDY *– AMAZON'S* *HIRING* *TOOL*

- Trained on 10 years of resumes (mostly male)

- Penalized resumes with “women’s” terms (e.g., “women’s chess club”)

- Discontinued after internal audit revealed bias

Should Amazon have caught this earlier?
What governance was missing?

DATA BIAS

- Occurs when training data is skewed or incomplete
- Underrepresentation of minorities, women, rural populations
- Example: Medical AI trained mostly on white male patients

If your data doesn't reflect the world, your AI won't either

SAMPLING BIAS

- Data collected from non-representative sources
- Example: Social media data skewed toward younger, urban users
- Consequence: AI fails to generalize to broader population

Sampling bias is like polling only one neighborhood and claiming to know the whole city

LABEL BIAS

- Human-labeled data may reflect subjective or biased judgments
- Example: Performance reviews influenced by gender stereotypes
- AI learns these biases as “truth”

Two identical resumes labeled differently

MEASUREMENT BIAS

- Using flawed proxies for real-world traits
- Example: Using postal code as a proxy for creditworthiness
- Can reinforce socioeconomic and racial disparities

Measurement bias is subtle but dangerous—it hides behind “business logic.”

ALGORITHMIC BIAS



- MODEL ARCHITECTURE OR TRAINING PROCESS INTRODUCES BIAS



- EXAMPLE: DECISION TREES FAVORING MAJORITY CLASS



- BIAS CAN EMERGE EVEN WITH "CLEAN" DATA

SOCIETAL BIAS

- AI reflects and amplifies existing inequalities
- Example: Predictive policing targeting minority neighborhoods
- Ethical dilemma: Should AI “mirror” society or improve it?

AI doesn't just reflect society—it can shape it.
That's a huge responsibility

FEEDBACK LOOPS

- Biased outcomes reinforce biased data
- Example: Loan denials reduce future credit data for marginalized groups
- Creates self-fulfilling cycles of exclusion
 - Circular loop “Bias → Outcome → Data → Bias”

BIAS IN TRAINING VS. DEPLOYMENT

- Bias can enter at multiple stages:
- Data collection
- Model training
- Deployment context
- Each stage requires different mitigation strategies

Bias isn't a bug—it's a lifecycle issue. You need governance at every step.

SUMMARY – SOURCES OF BIAS

- Bias is multifaceted: data, design, deployment
- Fairness requires intentional design and oversight
- Business leaders must ask:
- Who benefits?
- Who is harmed?
- Who is missing?

These questions aren't just ethical—they're strategic. They shape your brand, your culture, and your bottom line.

*MITIGATION
STRATEGIES &
LEADERSHIP
RESPONSIBILITY*



GROUP FAIRNESS VS. INDIVIDUAL FAIRNESS

- Group Fairness: Equal outcomes across demographic groups

Example: Same loan approval rate for men and women

- Individual Fairness: Similar individuals treated similarly

Example: Two applicants with identical credit profiles get same decision

- Trade-off: Group fairness may sacrifice individual nuance

Fairness isn't binary—it's a balancing act. Your business context determines which lens matters more.

COUNTERFACTUAL FAIRNESS

- Definition: A decision is fair if it would remain the same in a counterfactual world where the individual belonged to a different group
- Example: Would this person still be denied a loan if they were of a different race?
This is fairness through imagination—asking “what if” to uncover hidden bias.

DISPARATE IMPACT ANALYSIS

- Measures whether outcomes disproportionately affect protected groups
- Common in HR, lending, housing
- Legal standard in many jurisdictions

Disparate impact doesn't require intent—just unequal results. That's why it's so powerful

FAIRNESS METRICS

- Demographic Parity: Equal positive outcomes across groups
- Equal Opportunity: Equal true positive rates
- Calibration: Predictions are equally accurate across groups
- No single metric fits all—choose based on context

Metrics are tools, not truths. Use them to guide—not dictate—your fairness strategy.



*TRADE-
OFFS
BETWEEN
ACCURACY
& FAIRNESS*

-
- Improving fairness may reduce predictive accuracy
-
- Example: Adjusting model to reduce gender bias may lower precision
-
- Business leaders must weigh ethical vs. operational goals
-

Would you accept lower accuracy for higher fairness?

*DIVERSE &
REPRESENTATIVE
DATA*

- Include varied demographics, geographies, and contexts
- Avoid overfitting to dominant groups
- Use synthetic data to fill gaps if needed

Diversity in data is like diversity in teams—it leads to better decisions.

BIAS AUDITS & TESTING

- Regularly test models for bias across key metrics
- Use third-party audits for transparency
- Document findings and mitigation steps

Facebook's internal audit revealed ad targeting bias

Bias audits are your ethical smoke detectors—don't skip them.

EXPLAINABLE AI (XAI)

- Makes AI decisions interpretable to humans
- Helps detect bias and build trust
- Tools: SHAP, LIME, feature importance scores

Feature importance refers to techniques that calculate a score for all the input features for a given model. The scores represent the “importance” of each feature.

If you can't explain your AI, you can't defend it.
Transparency is non-negotiable.

HUMAN-IN- THE-LOOP SYSTEMS



- Combine AI with human oversight



- Humans validate or override AI decisions



- Useful in high-stakes areas: hiring, healthcare, finance



- AI should assist—not replace—human judgment in sensitive domains.

INCLUSIVE DESIGN PRACTICES

- Involve diverse stakeholders in design and testing

- Use personas from marginalized groups

- Conduct fairness reviews during development

Inclusion isn't just about who uses your product—it's about who builds it.

TRANSPARENCY & DOCUMENTATION

Model Card

Model Details

Basic information about the model.

- Person or organization developing model
- Model date
- Model version
- Model type
- Information about training algorithms, parameters, fairness constraints or other applied approaches, and features
- Paper or other resource for more information
- Citation details
- License
- Where to send questions or comments about the model

Intended Use

Use cases that were envisioned during development.

- Primary intended uses
- Out-of-scope use cases

Factors

Factors could include demographic or phenotypic groups, environmental conditions, technical attributes, or others listed as required.

- Relevant factors
- Evaluation factors

Metrics

Metrics should be chosen to reflect potential real-world impacts of the model.

- Model performance measures
- Decision thresholds
- Variation approaches

Evaluation Data

Details on the dataset(s) used for the quantitative analyses in the card.

- Datasets
- Motivation
- Preprocessing

Training Data

May not be possible to provide in practice. When possible, this section should mirror Evaluation Data. If such detail is not possible, minimal allowable information should be provided here, such as details of the distribution over various factors in the training datasets.

Quantitative Analyses

- Unitary results
- Intersectional results

Ethical Considerations

Caveats and Recommendations

- Maintain model cards and datasheets
- Document assumptions, limitations, and fairness goals
- Share with internal and external stakeholders
- Example:
- Google's Model Cards for AI transparency
 - Documentation isn't bureaucracy—it's accountability.



IMPACT OF BIAS ON BUSINESS

Reputational Damage

Bias in AI can expose discriminatory practices, harming a company's reputation and public image.

Legal and Compliance Risks

Non-compliance with anti-discrimination and data laws creates significant legal risks for businesses.

Customer Trust and Loyalty

Loss of trust from biased systems can reduce brand loyalty and market share significantly.

Operational Inefficiency

Biased AI may cause poor decisions, reducing operational efficiency and effectiveness.

**QUALITY
CONTROL CUSTOMERS
STRATEGY TESTING
ANALYSIS BRAND
VERIFICATION**



REGULATORY LANDSCAPE AND COMPLIANCE

Global AI Regulation Trends

Governments worldwide are creating AI regulations focusing on bias, fairness, and ethical governance.

EU AI Act Requirements

The EU AI Act mandates transparency, human oversight, and fairness for high-risk AI systems.

Compliance Strategies

Organizations must implement ethical audits, documentation, and stakeholder engagement to ensure compliance.

Benefits of Compliance

Proactive compliance builds public trust and offers a competitive advantage in the AI market.

ETHICAL AI GOVERNANCE MODELS

- Create cross-functional ethics committees
- Define escalation paths for bias incidents
- Align AI strategy with corporate values

Governance turns good intentions into sustainable practice.

01

Data Ethics

Implement data privacy standards to protect user information.

02

Compliance Frameworks

Integrate AI systems with existing legal regulations and guidelines.

03

Risk Assessment

Regularly evaluate AI systems for potential ethical risks and biases.

04

Transparency Measures

Enhance algorithmic transparency through explainable AI practices.

05

Stakeholder Engagement

Involve diverse stakeholders in AI development and decision-making processes.

06

Continuous Monitoring

Establish ongoing oversight for AI systems post-deployment.

AI Governance and Regulatory Frameworks

BUILDING DIVERSE AI TEAMS

- Diversity in teams reduces blind spots
- Include ethicists, sociologists, and domain experts
- Foster psychological safety for dissenting views

Homogeneous teams build homogeneous AI.
Diversity isn't optional—it's foundational.

STAKEHOLDER ENGAGEMENT

- Engage affected communities in design and testing
- Use participatory methods: surveys, focus groups, co-design
- Build trust through dialogue, not just data

Fairness isn't something you do to people—it's something you build with them.

CREATING AN AI ETHICS BOARD

- Internal or external advisory group
- Reviews high-risk models and decisions
- Provides ethical guidance and accountability

Example:

- Salesforce's Office of Ethical and Humane Use of Technology

WHAT KIND OF LEADER WILL YOU BE?

- Will you prioritize fairness even when it's inconvenient?
- Will you challenge biased systems?
- Will you build teams that reflect the world?

Leadership in AI isn't about knowing the answers—it's about asking the right questions.

YOUR ROLE IN SHAPING FAIR AI

- Fairness is a strategic, ethical, and human imperative
- Bias mitigation requires technical, organizational, and cultural shifts
- Business leaders must lead with courage, clarity, and compassion

You're not just shaping products—you're shaping society. Let's do it right.

RECAP OF KEY CONCEPTS

- Bias in AI: Systematic unfairness from data, design, or deployment
- Fairness: Equitable treatment across individuals or groups
- Business Impact: Legal, reputational, financial risks
- Mitigation: Diverse data, audits, transparency, inclusive design
- Leadership: Ethical governance and stakeholder engagement

We've covered the what, why, and how. Now let's talk about you—your role in shaping the future.

FINAL THOUGHTS – THE FUTURE OF ETHICAL AI

- AI is not neutral—it reflects human values
- Fairness is a moving target, not a checkbox
- Business leaders must lead with empathy, integrity, and foresight
- Your decisions shape not just markets—but lives

You're not just future executives—you're future stewards of technology. Let's make it just.

*“TECHNOLOGY IS A MIRROR—IT REFLECTS
WHO WE ARE. LET’S MAKE SURE IT REFLECTS
OUR BEST”*