## Reflection on Financial Services: Making AI Decisions Reviewable

The financial-services lecture material makes me think about AI adoption as a problem of responsibility. A system can produce a useful answer quickly while leaving the organization unable to explain who approved it or which evidence supported it. The speaker's questions about checking, sign-off and later proof are therefore more useful to me than a general claim that AI improves productivity. A bank needs a way to reconstruct a consequential decision after the immediate conversation has ended.

The discussion of human review exposes a practical difficulty. If an employee must repeat every step to verify an AI-generated result, much of the time saving disappears. Yet replacing that employee with another model does not automatically produce independent evidence. A generator and an AI reviewer could both rely on the same outdated document. I would separate checks that can be defined precisely, such as verifying a source version or recalculating a total, from judgments that need an accountable person. The relevant measure is the quality and cost of the whole process, including review and correction.

The project-management demonstration gives this principle a concrete form. Meeting text becomes linked actions, dates and views, but conflicting dates are flagged rather than silently resolved. I find this more convincing than a polished output that hides uncertainty. Keeping a task connected to its source allows a person to examine the proposed change. However, a source link only shows where a statement came from; it does not prove that the statement is correct or that the person who made it had authority. Important changes still need an explicit owner and approval.

The distinction between productivity support and a core business process also changes how I would assess a use case. The same dashboard may begin as an exploratory aid and later become an input to recurring decisions. Its software interface has not changed, but the consequences of an error have. I would reassess access, checking and escalation when the use changes, instead of assuming that the original approval covers every later purpose.

My main takeaway is to start with a bounded workflow whose evidence and failure paths can be inspected. I would evaluate a pilot using representative normal and difficult cases, record missed errors as well as unnecessary escalations, and include human review time in the cost. These are proposed evaluation steps, not results demonstrated in the material. They would help decide whether automation has improved the work sufficiently to justify expanding it.

Source: Financial-services lecture recording, source A (S26), 08:23-09:21, 25:37-26:38, 44:18-48:53 and 55:38-59:24. Available recording excerpts are incomplete; the demonstrations are not evidence of a production deployment.

I used OpenAI Codex to assist with this reflection.